

# Cardboard Box Detection and Localization using RetinaNet (Keras)

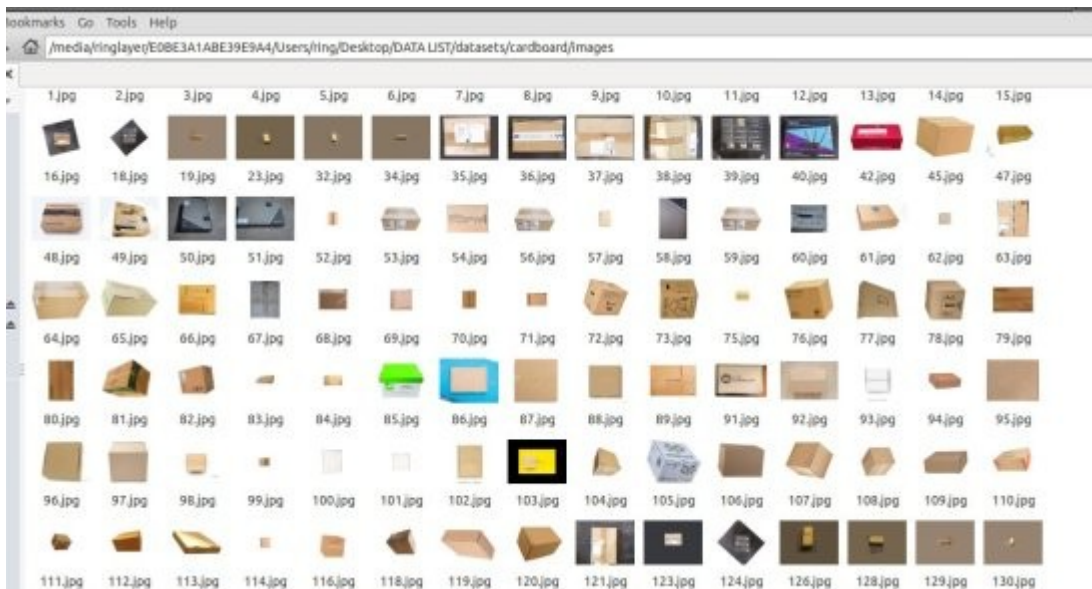
Antonius Robotsoft - [www.robotsoft.co.id](http://www.robotsoft.co.id) – <https://github.com/antoniusrobotsoft>

Keras RetinaNet is keras implementation of RetinaNet object detection as described in Focal Loss for Dense Object Detection Paper by Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, Piotr Dollar (<https://arxiv.org/abs/1708.02002>).

We are going to train a model using keras retinanet to detect and localize a custom object : “cardboard box” from the image.

## Step 1. Dataset Preparation

In order to detect cardboard boxes in the image, I have prepared several cardboard box images :



The dataset can be downloaded from :

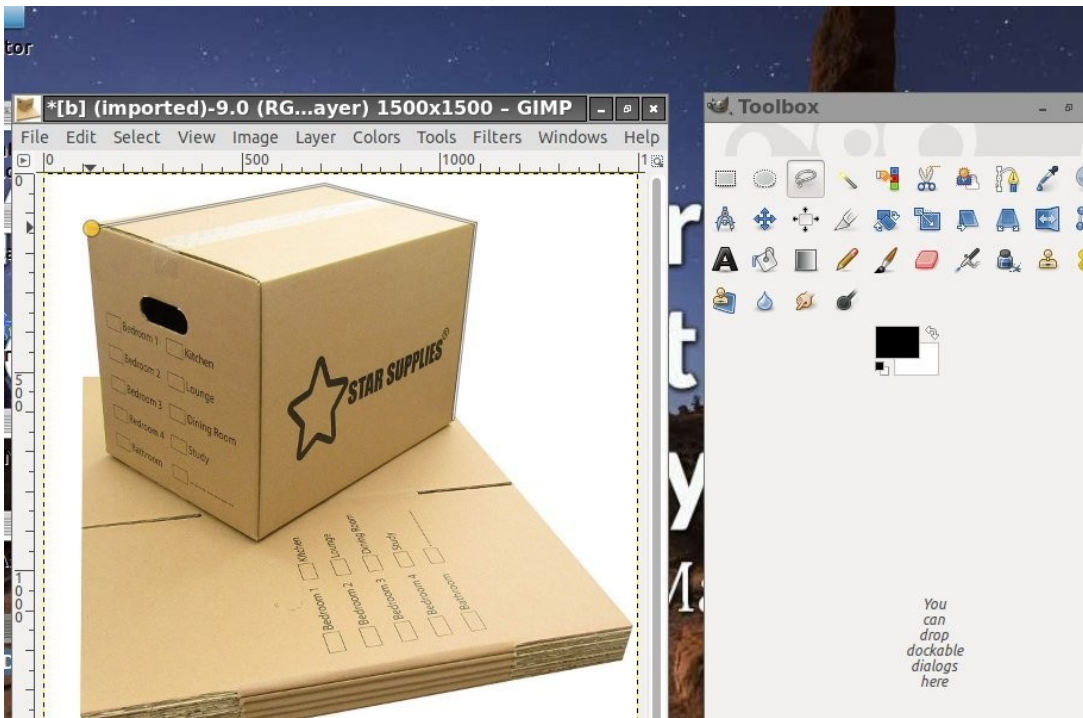
<https://jfid.sjasaplus.com/repo/cardboard.tar.bz2>

<https://jfid.sjasaplus.com/repo/cardboard2.tar.bz2>

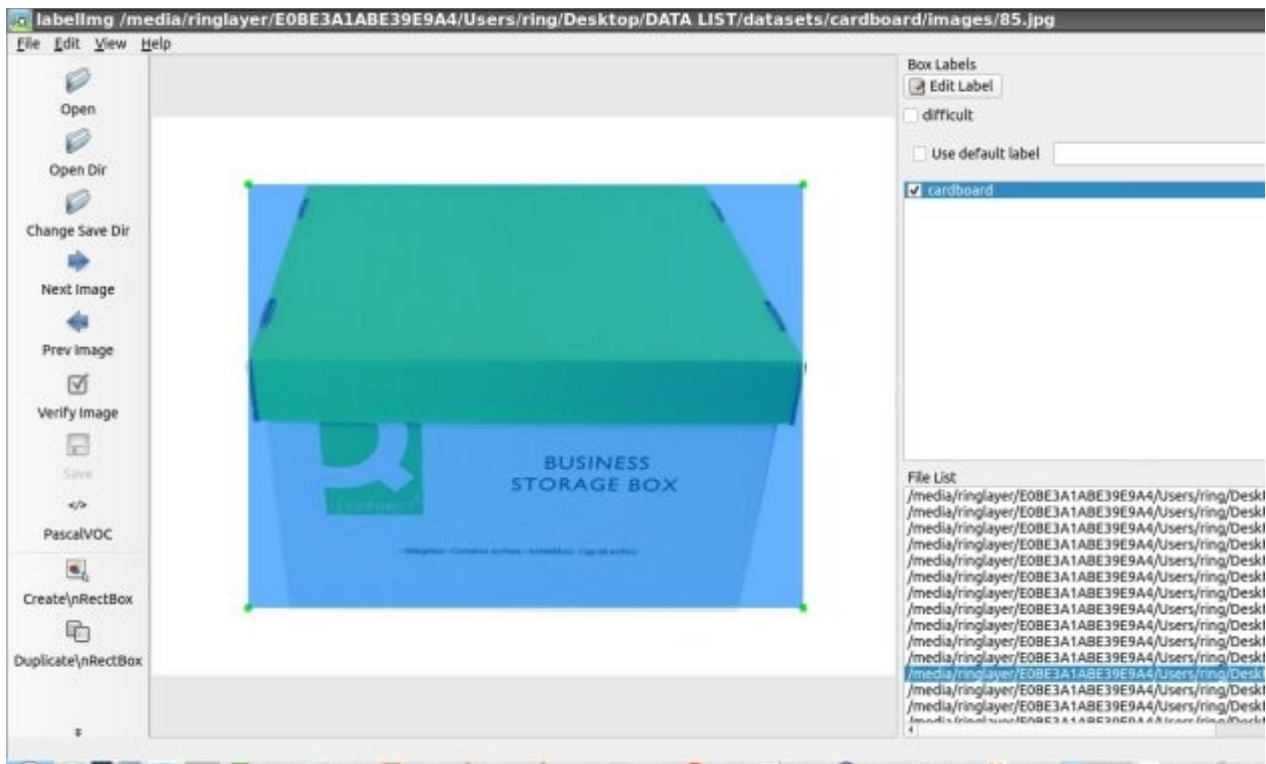
Keras Retinanet will resize any input image before input layer, the input image will be resized into an image with the maximal height 800px. For example, when we give input image 640×480 it will be resized into 1067×800. When we give input image 800×600, it will be resized into 1067×800.

This dataset was created using gimp and labelImg. It contains 640×480 images. We are going to train a new model using this dataset to detect image from webcam with 640×480 resolution.

Gimp was used to extract only relevant part of the image using lasso tool :



labelling was used for labelling specific areas within image that contains our object. Do not use too much margin when labelling the object since the quality of our detection depends on our dataset quality.



## Step 2. Converting the Dataset from Pascal VOC Format into CSV Annotation Format for Keras RetinaNet

Next, we need to convert the dataset from voc into csv. We are going to use `conv_pascal_to_csv`, this utility can be cloned from github url :

[https://github.com/akchan/conv\\_pascal\\_to\\_csv.git](https://github.com/akchan/conv_pascal_to_csv.git)

```
$ ruby conv_pascal_to_csv.rb -help
```

```
Usage: conv_pascal_to_csv [options]
```

```
-annotation-path PATH path to xml annotations directory
```

```
-image-path PATH path to jpeg images directory
```

```
-val-ratio float sample ratio for validation (0.0 – 1.0). default=0.1
```

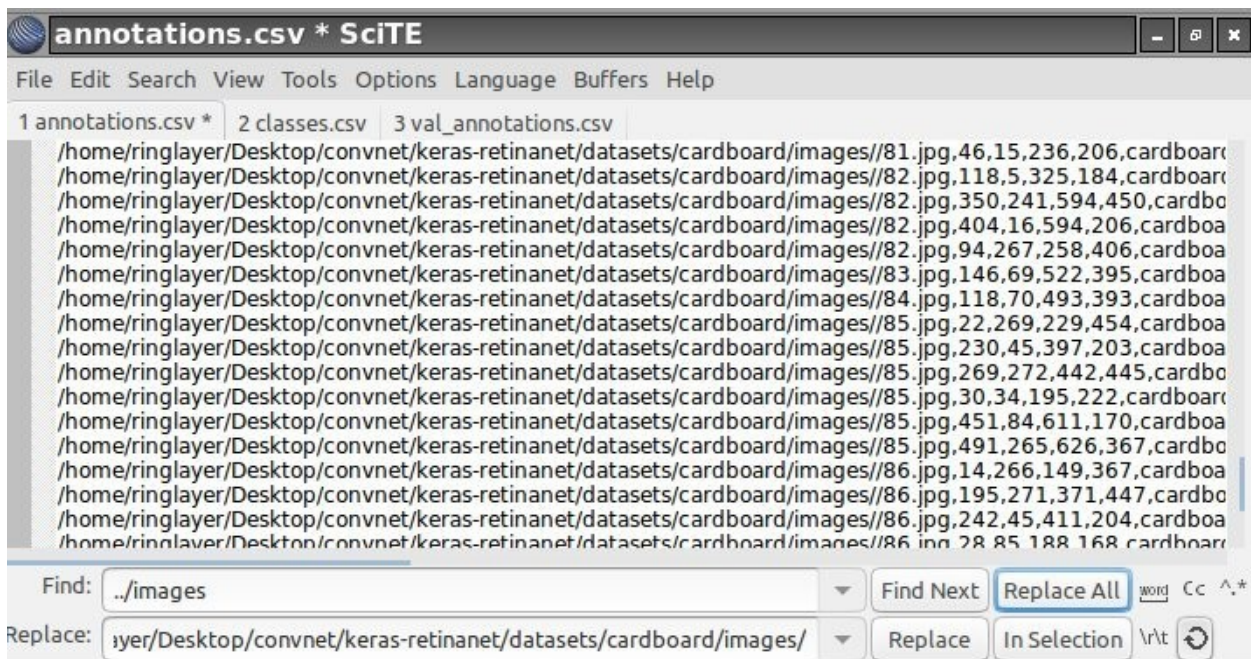
```
$ ruby conv_pascal_to_csv.rb -annotation-path /home/ringlayer/Desktop/convnet/keras-retinanet/datasets/cardboard2/voclabels -image-path /home/ringlayer/Desktop/convnet/keras-retinanet/datasets/cardboard2/images
```

Once the conversion finished, we will get 3 files in csv directory :

```
$ ls csv
```

```
annotations.csv classes.csv val_annotations.csv
```

I replaced image paths at `annotations.csv` and `val_annotations.csv` with absolute paths of my dataset location :



### Step 3. Start Training New Model

In order to train new model, we can do **transfer learning** by using previously trained weight. Here, I use Resnet50 as backbone to train new model :

```
python3 keras_retinanet/bin/train.py --freeze-backbone --random-transform --weights
weights/ResNet-50-model.keras.h5 --batch-size 8 --steps 500 --epochs 15 csv csv/annotations.csv
csv/classes.csv
```

These parameters required since we were not feeding the neural net at once with all training images, instead we will divide it into several epochs (each epoch will be divided into 500 steps)

#### –batch-size 8

We used batch size 8, The batch size defines the number of images that will be propagated through the network.

#### –epochs 15

One epoch means an entire dataset is passed forward and backward through the neural network. Since we used 15 epochs, it means it will need 15 epochs to complete the entire dataset provided at annotations.csv.

#### –steps 500

We use 500 steps on each epoch. This is the number of batch iterations before a training epoch is considered finished.

Here, I trained the model using gtx1080 with 8 gb vram :



```

ringlayer@ringlayer-CORSAIR-ONE:~/Desktop/convnet/keras-retinanet$ ./train
Using TensorFlow backend.
2019-07-02 01:30:34.978905: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:964] successful NUMA node read
t least one NUMA node, so returning NUMA node zero
2019-07-02 01:30:34.979291: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1432] Found device 0 with properti
name: GeForce GTX 1080 major: 6 minor: 1 memoryClockRate(GHz): 1.7335
pciBusID: 0000:01:00.0
totalMemory: 7.93GiB freeMemory: 7.75GiB
2019-07-02 01:30:34.979306: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1511] Adding visible gpu devices:
2019-07-02 01:30:35.155345: I tensorflow/core/common_runtime/gpu/gpu_device.cc:982] Device interconnect StreamExe
2019-07-02 01:30:35.155374: I tensorflow/core/common_runtime/gpu/gpu_device.cc:988] 0
2019-07-02 01:30:35.155383: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1001] 0: N
2019-07-02 01:30:35.155506: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1115] Created TensorFlow device (,
memory) -> physical GPU (device: 0, name: GeForce GTX 1080, pci bus id: 0000:01:00.0, compute capability: 6.1)
2019-07-02 01:30:35.155983: I tensorflow/core/common_runtime/process_util.cc:69] Creating new thread pool with de
sm threads for best performance.
path : csv/classes.csv
path : csv/annotations.csv
Creating model, this may take a second...

```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, None, None, 3 0		
conv1 (Conv2D)	(None, None, None, 6 9408		input_1[0][0]
bn_conv1 (BatchNormalization)	(None, None, None, 6 256		conv1[0][0]
conv1_relu (Activation)	(None, None, None, 6 0		bn_conv1[0][0]
pool1 (MaxPooling2D)	(None, None, None, 6 0		conv1_relu[0][0]
res2a_branch2a (Conv2D)	(None, None, None, 6 4096		pool1[0][0]
bn2a_branch2a (BatchNormalizati	(None, None, None, 6 256		res2a_branch2a[0][0]
res2a_branch2a_relu (Activation	(None, None, None, 6 0		bn2a_branch2a[0][0]

This configuration takes about 7gb vram

```

ringlayer@ringlayer-CORSAIR-ONE:~/Desktop/convnet/darknet$ nvidia-smi
Tue Jul  2 01:47:04 2019
+-----+
| NVIDIA-SMI 410.48                  Driver Version: 410.48          |
+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+
|  0  GeForce GTX 1080    Off      | 00000000:01:00.0  On  |          N/A         |
| 54%   66C    P2      165W / 180W | 7845MiB / 8116MiB |    99%    Default   |
+-----+-----+

+-----+
| Processes:                         GPU Memory Usage |
| GPU       PID    Type    Process name                     | Memory Usage |
+-----+-----+
|    0      1135     G   /usr/lib/xorg/Xorg                     |    52MiB     |
|    0      6465     C   python3                             |   7781MiB    |
+-----+

```

In order to fine tune new model training we can use a gpu with 11 gb vram.

```

ringlayer@ringlayer-CORSAIR-ONE: ~
File Edit Tabs Help
ringlayer@ri... X ringlayer@ri... X
500/500 [=====] - 429s 857ms/step - loss: 0.6098 - regression_loss: 0.5254 - classification_loss: 0.0844
Epoch 00008: saving model to ./snapshots/resnet50_csv_08.h5
Epoch 9/15
500/500 [=====] - 429s 857ms/step - loss: 0.5562 - regression_loss: 0.4848 - classification_loss: 0.0714
Epoch 00009: saving model to ./snapshots/resnet50_csv_09.h5
Epoch 10/15
500/500 [=====] - 429s 858ms/step - loss: 0.5186 - regression_loss: 0.4550 - classification_loss: 0.0636
Epoch 00010: saving model to ./snapshots/resnet50_csv_10.h5
Epoch 11/15
500/500 [=====] - 429s 858ms/step - loss: 0.4861 - regression_loss: 0.4297 - classification_loss: 0.0564
Epoch 00011: saving model to ./snapshots/resnet50_csv_11.h5
Epoch 12/15
500/500 [=====] - 429s 858ms/step - loss: 0.4627 - regression_loss: 0.4114 - classification_loss: 0.0513
Epoch 00012: saving model to ./snapshots/resnet50_csv_12.h5
Epoch 13/15
500/500 [=====] - 428s 857ms/step - loss: 0.4351 - regression_loss: 0.3903 - classification_loss: 0.0448
Epoch 00013: saving model to ./snapshots/resnet50_csv_13.h5
Epoch 14/15
500/500 [=====] - 428s 857ms/step - loss: 0.4162 - regression_loss: 0.3744 - classification_loss: 0.0418
Epoch 00014: saving model to ./snapshots/resnet50_csv_14.h5
Epoch 15/15
500/500 [=====] - 429s 857ms/step - loss: 0.3975 - regression_loss: 0.3598 - classification_loss: 0.0378
Epoch 00015: saving model to ./snapshots/resnet50_csv_15.h5
ringlayer@ringlayer-CORSAIR-ONE:~/Desktop/convnet/keras-retinanet$ cd

```

The trained weight on each epoch will be saved to snapshots directory. We are going to use resnet50\_csv\_15.h5 for our final model to detect cardboard (I renamed resnet50\_csv\_15.h5 into cb20.h5).

The source code for testing this model on webcam and video can be fetched from github :

<https://github.com/antoniuserobotsoft>

Testing our model to detect cardboard in video and image :

